

PML-LAB-2_maha18

March 15, 2021

0.0.1 Lab2. Pizza Liking Prediction using kNN

0.0.2 Step2. [Import dataset]. Using Pandas, import “pizza.csv” file and print properties such as head(), shape, columns and info.

```
[5]: import pandas as pd  
import csv
```

```
[6]: piz=pd.read_csv("pizza.csv")
```

```
[7]: piz.head()
```

```
[7]:
```

	age	weight	likepizza
0	50	65	0
1	20	55	1
2	15	40	1
3	70	65	0
4	30	70	1

```
[8]: piz.tail()
```

```
[8]:
```

	age	weight	likepizza
1	20	55	1
2	15	40	1
3	70	65	0
4	30	70	1
5	75	60	0

```
[9]: piz.shape
```

```
[9]: (6, 3)
```

```
[10]: df = pd.read_csv("pizza.csv")
```

```
[11]: df
```

```
[11]:
```

	age	weight	likepizza
0	50	65	0
1	20	55	1

2	15	40	1
3	70	65	0
4	30	70	1
5	75	60	0

```
[12]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6 entries, 0 to 5
Data columns (total 3 columns):
#   Column      Non-Null Count  Dtype
---  -
0    age         6 non-null      int64
1   weight      6 non-null      int64
2  likepizza    6 non-null      int64
dtypes: int64(3)
memory usage: 272.0 bytes
```

```
[13]: res = df.columns
```

```
[14]: print(res)
```


```
Index(['age', 'weight', 'likepizza'], dtype='object')
```

0.0.3 Step3. [Visualize Relationships]. Plot relplot between “age” and “weight”, with hue as “likePizza”

```
[15]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
```


```
[16]: sns.relplot(x="age", y="weight", hue="likepizza", data=piz)
```

```
[16]: <seaborn.axisgrid.FacetGrid at 0x20695e6b7c0>
```



```
[17]: sns.relplot(x='age', y='weight', data=piz, kind='scatter')
```

```
[17]: <seaborn.axisgrid.FacetGrid at 0x20695e38340>
```



```
[20]: y=piz.likepizza
```

0.0.4 Step4. [Prepare X matrix and y vector]. Extract “age” and “weight” columns and store into new dataframe X. Similarly, extract “likePizza” column and store into y.

```
[21]: dat = ['age','weight']  
X=piz[dat]
```

```
[22]: X
```

```
[22]:
```

	age	weight
0	50	65
1	20	55
2	15	40
3	70	65
4	30	70

```
5    75    60
```

```
[23]: y
```

```
[23]: 0    0
      1    1
      2    1
      3    0
      4    1
      5    0
      Name: likepizza, dtype: int64
```

0.0.5 Step5. [Examine X and y]. Print X, y, type of X and type of y.

```
[24]: X.dtypes
```

```
[24]: age      int64
      weight   int64
      dtype: object
```

```
[25]: y.dtype
```

```
[25]: dtype('int64')
```

0.0.6 Step6. [Model building]. Create KNeighborsClassifier(n_neighbors=2) from sklearn and perform fit on X and y.

```
[26]: !pip install sklearn
```

```
Requirement already satisfied: sklearn in c:\programdata\anaconda3\lib\site-
packages (0.0)
Requirement already satisfied: scikit-learn in
c:\programdata\anaconda3\lib\site-packages (from sklearn) (0.23.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in
c:\programdata\anaconda3\lib\site-packages (from scikit-learn->sklearn) (2.1.0)
Requirement already satisfied: scipy>=0.19.1 in
c:\programdata\anaconda3\lib\site-packages (from scikit-learn->sklearn) (1.5.2)
Requirement already satisfied: joblib>=0.11 in
c:\programdata\anaconda3\lib\site-packages (from scikit-learn->sklearn) (0.17.0)
Requirement already satisfied: numpy>=1.13.3 in
c:\programdata\anaconda3\lib\site-packages (from scikit-learn->sklearn) (1.19.2)
```

```
[27]: from sklearn.neighbors import KNeighborsClassifier
```

```
[28]: pizza_piz = KNeighborsClassifier(n_neighbors=2)
      pizza_piz.fit(X, y)
```

```
[28]: KNeighborsClassifier(n_neighbors=2)
```

0.0.7 Step7. [Model testing]. Using your KNN model, predict if a person will like Pizza or not.

```
[29]: print(pizza_piz.predict(X))
```

```
[0 1 1 0 1 0]
```

0.0.8 Will a person who is 25 years with weight 50 kgs like Pizza or not? – The answer should be 1 (ie., YES)

```
[30]: new=[25,50]
      pizza_piz.predict([new])
```

```
[30]: array([1], dtype=int64)
```

0.0.9 Will a person who is 60 years with weight 60 kgs like Pizza or not? – The answer should be 0 (ie., NO)

```
[31]: new=[60,60]
      pizza_piz.predict([new])
```

```
[31]: array([0], dtype=int64)
```

0.0.10 Step.8 [Change n_neighbors = 3]. Now, create new model, perform fit and predict steps. Check results for the above 2 queries. Are they same?

```
[32]: pizza_piz = KNeighborsClassifier(n_neighbors=3)
      pizza_piz.fit(X, y)
```

```
[32]: KNeighborsClassifier(n_neighbors=3)
```

```
[33]: print(pizza_piz.predict(X))
```

```
[0 1 1 0 1 0]
```

```
[34]: new=[25,50]
      pizza_piz.predict([new])
```

```
[34]: array([1], dtype=int64)
```

```
[35]: new=[60,60]
      pizza_piz.predict([new])
```

```
[35]: array([0], dtype=int64)
```

0.0.11 Step9. [Predict on entire dataset]. Now, perform prediction on entire X matrix and store result as y_pred.

```
[36]: y_pred=pizza_piz.predict(X)
```

```
[37]: y_pred
```

```
[37]: array([0, 1, 1, 0, 1, 0], dtype=int64)
```

0.0.12 Step10. [Accuracy function]. Create a function accuracy() and returns accuracy.

```
[74]: def accuracy(actual,pred):  
      return sum(actual == pred) / float(actual.shape[0])
```

0.0.13 Step11. [Find accuracy]. Call accuracy() with y and y_pred as parameters and print accuracy score. Are you getting score as 1.0 ?

```
[76]: accuracy(y,y_pred)
```

```
[76]: 0.5
```

0.0.14 Step12. [Prediction on Test Set]

Using Pandas, import “pizza_test.csv” file and print properties such as head(), shape, columns and info. Using KNN model with n_neighbors=2, that you created previously, perform prediction on X values from pizza_test dataframe. Call accuracy function and print accuracy score. Are you getting a score of 0.5? That is, our model has predicted 2 samples correctly and two wrongly, out of 4 samples in the test set.

```
[38]: pizz=pd.read_csv("pizza_test.csv")
```

```
[39]: pizz.head()
```

```
[39]:   age  weight  likepizza  
0   48     68         1  
1   35     45         1  
2   15     40         0  
3   55     65         0
```

```
[40]: pizz.tail()
```

```
[40]:   age  weight  likepizza  
0   48     68         1  
1   35     45         1  
2   15     40         0  
3   55     65         0
```

```
[41]: pizz.shape
```

```
[41]: (4, 3)
```

```
[42]: df = pd.read_csv("pizza_test.csv")
```

```
[43]: df
```

```
[43]:   age  weight  likepizza
0   48     68         1
1   35     45         1
2   15     40         0
3   55     65         0
```

```
[44]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4 entries, 0 to 3
Data columns (total 3 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         4 non-null      int64
1   weight      4 non-null      int64
2   likepizza   4 non-null      int64
dtypes: int64(3)
memory usage: 224.0 bytes
```


```
[45]: res = df.columns
```

```
[46]: print(res)
```

```
Index(['age', 'weight', 'likepizza'], dtype='object')
```


```
[47]: sns.relplot(x="age", y="weight", hue="likepizza", data=pizz)
```

```
[47]: <seaborn.axisgrid.FacetGrid at 0x20697feb880>
```

```
[48]: sns.relplot(x='age', y='weight', data=pizz, kind='scatter')
```

```
[48]: <seaborn.axisgrid.FacetGrid at 0x206980419a0>
```



```
[49]: rk = pizz.dropna(axis=0)
```

```
[50]: rk
```

```
[50]:
```

	age	weight	likepizza
0	48	68	1
1	35	45	1
2	15	40	0
3	55	65	0

```
[51]: y=pizz.likepizza
```

```
[52]: datt = ['age','weight']  
X=pizz[datt]
```

```
[53]: X
```

```
[53]:   age  weight
      0   48     68
      1   35     45
      2   15     40
      3   55     65
```

```
[54]: y
```

```
[54]: 0    1
      1    1
      2    0
      3    0
      Name: likepizza, dtype: int64
```

```
[55]: X.dtypes
```

```
[55]: age      int64
      weight  int64
      dtype: object
```

```
[56]: y.dtypes
```

```
[56]: dtype('int64')
```

```
[57]: from sklearn.neighbors import KNeighborsClassifier
```

```
[81]: piz1 = KNeighborsClassifier(n_neighbors=2)
      piz1.fit(X, y)
```

```
[81]: KNeighborsClassifier(n_neighbors=2)
```

```
[83]: (piz1.predict(X))
```

```
[83]: array([0, 0, 0, 0], dtype=int64)
```

```
[84]: def accuracy(actual, pred):
      return sum(actual == pred) / float(actual.shape[0])
```

```
[85]: y_pred=piz1.predict(X)
```

```
[86]: accuracy(y,y_pred)
```

```
[86]: 0.5
```

0.0.15 step13. [Find best value for k]. If you want to improve the accuracy of your model, then you should use the best value k for the nearest neighbors.

```
[77]: score = []  
      for k in range(1,4):  
          best = KNeighborsClassifier(n_neighbors=k)  
          best.fit(X, y)  
          best.predict(X)  
          y_predt = best.predict(X)  
          acc=accuracy(y,y_predt)  
          score.append((k,acc))
```

```
[78]: score
```

```
[78]: [(1, 1.0), (2, 0.5), (3, 0.5)]
```

0.0.16 Step14. [accuracy_score function]. Call accuracy_score() function with y_test and y_pred values. You can import as “from sklearn.metrics import accuracy_score”.

```
[88]: from sklearn.metrics import accuracy_score
```

```
[89]: accuracy_score(y,y_predt)
```

```
[89]: 0.5
```