

PML_lab-3_205229118_Mahalakshmi.S

March 22, 2021

0.0.1 Lab3. Fuel Amount Prediction using Linear Regression

0.0.2 Step1. [Prepare your dataset]. Create fuel_data.csv file as shown above.

```
[2]: import pandas as pd
import csv
```

```
[3]: fuel=pd.read_csv("fuel_data.csv")
fuel
```

```
[3]:
```

	drivenKM	fuelAmount
0	390.00	3600.0
1	403.00	3705.0
2	396.50	3471.0
3	383.50	3250.5
4	321.10	3263.7
5	391.30	3445.2
6	386.10	3679.0
7	371.80	3744.5
8	404.30	3809.0
9	392.20	3905.0
10	386.43	3874.0
11	395.20	3910.0
12	381.00	4020.7
13	372.00	3622.0
14	397.00	3450.5
15	407.00	4179.0
16	372.40	3454.2
17	375.60	3883.8
18	399.00	4235.9

0.0.3 Step2. [Import dataset]. Using Pandas, import “fuel_data.csv” file and print properties such as head(), shape, columns, type and info.

```
[4]: fuel.head()
```

```
[4]:
```

	drivenKM	fuelAmount
0	390.0	3600.0

1	403.0	3705.0
2	396.5	3471.0
3	383.5	3250.5
4	321.1	3263.7

```
[5]: fuel.tail()
```

```
[5]:      drivenKM  fuelAmount
14      397.0      3450.5
15      407.0      4179.0
16      372.4      3454.2
17      375.6      3883.8
18      399.0      4235.9
```

```
[6]: fuel.shape
```

```
[6]: (19, 2)
```

```
[7]: df = pd.read_csv("fuel_data.csv")
```

```
[8]: df
```

```
[8]:      drivenKM  fuelAmount
0      390.00      3600.0
1      403.00      3705.0
2      396.50      3471.0
3      383.50      3250.5
4      321.10      3263.7
5      391.30      3445.2
6      386.10      3679.0
7      371.80      3744.5
8      404.30      3809.0
9      392.20      3905.0
10     386.43      3874.0
11     395.20      3910.0
12     381.00      4020.7
13     372.00      3622.0
14     397.00      3450.5
15     407.00      4179.0
16     372.40      3454.2
17     375.60      3883.8
18     399.00      4235.9
```

```
[9]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 19 entries, 0 to 18
Data columns (total 2 columns):
```

#	Column	Non-Null Count	Dtype
0	drivenKM	19 non-null	float64
1	fuelAmount	19 non-null	float64

dtypes: float64(2)
memory usage: 432.0 bytes

```
[10]: f = df.columns
      f
```

```
[10]: Index(['drivenKM', 'fuelAmount'], dtype='object')
```

0.0.4 Step3. [Preprocessing]. Check for missing values (Use isnull() method)

```
[11]: df.isnull()
```


```
[11]:   drivenKM  fuelAmount
0      False      False
1      False      False
2      False      False
3      False      False
4      False      False
5      False      False
6      False      False
7      False      False
8      False      False
9      False      False
10     False      False
11     False      False
12     False      False
13     False      False
14     False      False
15     False      False
16     False      False
17     False      False
18     False      False
```

0.0.5 Step4. [Visualize Relationships]. Plot relplot between “drivenKM” and “fuelAmount”.

```
[12]: import pandas as pd
      import matplotlib.pyplot as plt
      import seaborn as sns
      import numpy as np
```

```
[13]: sns.relplot(x="drivenKM", y="fuelAmount", data=fuel)
```

```
[13]: <seaborn.axisgrid.FacetGrid at 0x1e0f1afa7f0>
```



0.0.6 Step5. [Prepare X matrix and y vector]. Extract “drivenKM” column and store into new dataframe X. Similarly, extract “fuelAmount” and store into y.

```
[14]: data1 = ['drivenKM']
      X=fuel[data1]
      data2 = ['fuelAmount']
      y=fuel.fuelAmount
```

0.0.7 Step6. [Examine X and y]. Print X, y, type of X and type of y.

```
[15]: print(X)
      X.dtypes
```

```
      drivenKM
0      390.00
1      403.00
2      396.50
```

```
3      383.50
4      321.10
5      391.30
6      386.10
7      371.80
8      404.30
9      392.20
10     386.43
11     395.20
12     381.00
13     372.00
14     397.00
15     407.00
16     372.40
17     375.60
18     399.00
```

```
[15]: drivenKM    float64
      dtype: object
```

```
[16]: print(y)
      y.dtypes
```

```
0      3600.0
1      3705.0
2      3471.0
3      3250.5
4      3263.7
5      3445.2
6      3679.0
7      3744.5
8      3809.0
9      3905.0
10     3874.0
11     3910.0
12     4020.7
13     3622.0
14     3450.5
15     4179.0
16     3454.2
17     3883.8
18     4235.9
```

```
Name: fuelAmount, dtype: float64
```

```
[16]: dtype('float64')
```

0.0.8 Step7. [Split dataset]. Split dataset into 4 parts using `train_test_split()` method, such as `X_train`, `X_test`, `y_train` and `y_test`. Use 20% for test size. Later you can play around with this test size. Print the shape of all 4 parts.

```
[18]: from sklearn.model_selection import train_test_split
```

```
[19]: X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=0.  
      ↪8, test_size=0.2)
```

```
[20]: X_train, X_test, y_train, y_test
```

```
[20]: (   drivenKM  
      11    395.20  
       9    392.20  
       6    386.10  
       5    391.30  
      15    407.00  
       8    404.30  
      18    399.00  
      13    372.00  
      12    381.00  
       1    403.00  
      10    386.43  
      17    375.60  
       7    371.80  
      16    372.40  
       3    383.50,  
      drivenKM  
      14    397.0  
       0    390.0  
       4    321.1  
       2    396.5,  
      11   3910.0  
       9   3905.0  
       6   3679.0  
       5   3445.2  
      15   4179.0  
       8   3809.0  
      18   4235.9  
      13   3622.0  
      12   4020.7  
       1   3705.0  
      10   3874.0  
      17   3883.8  
       7   3744.5  
      16   3454.2  
       3   3250.5  
      Name: fuelAmount, dtype: float64,
```

```
14    3450.5
0     3600.0
4     3263.7
2     3471.0
Name: fuelAmount, dtype: float64)
```

```
[21]: X_train.shape
```

```
[21]: (15, 1)
```

```
[22]: X_test.shape
```

```
[22]: (4, 1)
```

```
[23]: y_train.shape
```

```
[23]: (15,)
```

```
[24]: y_test.shape
```

```
[24]: (4,)
```

0.0.9 Part-I. Linear Regression Baseline Model

0.0.10 Step8. [Build Model]. Create Linear Regression model and train with fit() using X_train and y_train values.

```
[25]: from sklearn.linear_model import LinearRegression
```

```
[26]: #create a linear regression object
model = LinearRegression()
#train a model
model.fit( X_train,y_train)
```

```
[26]: LinearRegression()
```

0.0.11 Step9. [Predict price for 800 KM]. If I need to travel 800 KM, how much do I need to spend on Diesel?. Are you getting this output, array([6905.64571567]).?

```
[30]: n=[[800]]
m=model.predict(n)
m
```

```
[30]: array([7799.94281895])
```

0.0.12 Step10. [Predict on entire dataset]. Now, perform prediction using entire `X_test` and store result as `y_pred`.

```
[33]: y_pred=model.predict(X_test)
      y_pred
```

```
[33]: array([3868.44701279, 3800.15800375, 3127.99904334, 3863.56922643])
```

0.0.13 Step11. [Print Mean Squared Error and R2 Error]. Are you getting output “MSE: 46181.0”. Also, print values of model parameters: `coef_` and `intercept_` values.

```
[37]: from sklearn.metrics import mean_squared_error
      from sklearn.metrics import r2_score
```

```
[38]: mse_ln=mean_squared_error(y_test,y_pred)
      mse_ln
```

```
[38]: 96817.06978545067
```

```
[39]: r2_score(y_test,y_pred)
```

```
[39]: -5.724087036950648
```

```
[40]: model.coef_
```

```
[40]: array([9.75557272])
```

```
[41]: model.intercept_
```

```
[41]: -4.515357046707777
```

0.0.14 Part-II. Linear Regression with Scaling using StandardScaler

0.0.15 Step12. [Normalize `X_train` and `X_test` values]. Use `StandardScaler`, scale `X_train` using `fit_transform()` method and `X_test` using `transform()` method.

```
[42]: from sklearn.preprocessing import StandardScaler
```

```
[43]: scaler = StandardScaler()
      ss3=scaler.fit_transform(X_train)    # scale X_train using fit_transform()
      ↪method
      print(ss3)
```

```
[[ 0.61206548]
 [ 0.35506309]
 [-0.16750845]
 [ 0.27796237]
 [ 1.62294157]]
```



```
[ 1.39163942]
[ 0.93760185]
[-1.37541971]
[-0.60441252]
[ 1.28027171]
[-0.13923819]
[-1.06701683]
[-1.3925532 ]
[-1.34115272]
[-0.39024386]]
```

```
[44]: ss5=scaler.transform(X_test)      #X_test using transform() method
      print(ss5)
```

```
[[ 0.76626692]
 [ 0.16659466]
 [-5.73589369]
 [ 0.72343319]]
```

0.0.16 Step13. [Build LR model]. Create a new LR model, fit on scaled X_train and predict on scaled X_test.

```
[45]: model1 = LinearRegression()
      model1.fit(ss3,y_train)
```

```
[45]: LinearRegression()
```

```
[46]: s1_y_pred = model1.predict(ss5)
      s1_y_pred
```

```
[46]: array([3868.44701279, 3800.15800375, 3127.99904334, 3863.56922643])
```

0.0.17 Step14. [Print Mean Squared Error and R2 Error]. What is the output?. MSE reduced or not?. Why?.

```
[48]: mean_squared_error(y_test,s1_y_pred)
```

```
[48]: 96817.06978545096
```


```
[49]: r2_score(y_test, s1_y_pred)
```

```
[49]: -5.724087036950668
```

0.0.18 Step15. [Plot scatter plot]. Display Scatter Plot between actual y (aka ground truth) vs predicted y values. That is, between y_test and y_pred.

```
[53]: plt.scatter(y_test,y_pred)
```

[53]: <matplotlib.collections.PathCollection at 0x1e0f52aa700>



0.0.19 Part-III. Linear Regression with Scaling using MinMaxScaler and Comparison with KNeighborsRegressor and SGDRegressor

0.0.20 Step16. [Repeat with MinmaxScaler]. Repeat scaling using MinMaxScaler, LR model creation, fit, predict and error computation steps.

```
[57]: from sklearn.preprocessing import MinMaxScaler  
mm_scaler = MinMaxScaler()
```

```
[58]: mm_ss = mm_scaler.fit_transform(X_train)  
mm_ss
```

```
[58]: array([[0.66477273],  
          [0.57954545],  
          [0.40625   ],  
          [0.55397727],  
          [1.        ],  
          [0.92329545],  
          [0.77272727],  
          [0.00568182],  
          [0.26136364],  
          [0.88636364],  
          [0.415625  ],  
          [0.10795455],
```

```
[0.          ],
 [0.01704545],
 [0.33238636]])
```

```
[59]: mm_ss5 = mm_scaler.transform(X_test)
mm_ss5
```

```
[59]: array([[ 0.71590909],
 [ 0.51704545],
 [-1.44034091],
 [ 0.70170455]])
```

```
[60]: model2 = LinearRegression()
model2.fit(mm_ss,y_train)
```

```
[60]: LinearRegression()
```

```
[62]: mms_y_pred = model2.predict(mm_ss5)
mms_y_pred
```

```
[62]: array([3868.44701279, 3800.15800375, 3127.99904334, 3863.56922643])
```

```
[63]: mean_squared_error(y_test,mms_y_pred)
```

```
[63]: 96817.06978545104
```

```
[64]: r2_score(y_test,mms_y_pred)
```

```
[64]: -5.724087036950673
```

0.0.21 Step17. [Compare KNN Regressor]. Repeat the above steps for KNeighborsRegressor model and compare MSE of LR with KNN Regressor.

```
[65]: from sklearn.neighbors import KNeighborsRegressor
```

```
[66]: m_neig = KNeighborsRegressor(n_neighbors=5)
m_neig.fit(X, y)
```

```
[66]: KNeighborsRegressor()
```

```
[67]: n1_y_pred = m_neig.predict(X)
n1_y_pred
```

```
[67]: array([3700.64, 3875.88, 3794.48, 3684.84, 3593.64, 3746.84, 3684.84,
 3745.04, 3875.88, 3666.24, 3569.74, 3794.48, 3741.6 , 3745.04,
 3794.48, 3875.88, 3745.04, 3745.04, 3754.48])
```

```
[68]: mse=mean_squared_error(y,n1_y_pred)
mse
```

```
[68]: 70460.30507368421
```

```
[69]: r2_score(y,n1_y_pred)
```

```
[69]: 0.06403925984775638
```

0.0.22 Step18. [Compare SGD Regressor]. Repeat the above steps for SGDRegressor model and compare MSE of LR with SGD Regressor.

```
[70]: from sklearn.linear_model import SGDRegressor
from sklearn.pipeline import make_pipeline
```

```
[71]: r = make_pipeline(StandardScaler(), SGDRegressor(max_iter=1000, tol=1e-3))
r.fit(X, y)
```

```
[71]: Pipeline(steps=[('standardscaler', StandardScaler()),
                    ('sgdregressor', SGDRegressor())])
```

```
[72]: re_y_pred = r.predict(X)
re_y_pred
```

```
[72]: array([3740.86178879, 3830.59155467, 3785.72667173, 3695.99690584,
          3265.29402959, 3749.83476537, 3713.94285902, 3615.24011655,
          3839.56453126, 3756.04682609, 3716.22061462, 3776.75369514,
          3678.74118163, 3616.62057448, 3789.17781657, 3858.20071341,
          3619.38149036, 3641.46881734, 3802.98239594])
```

```
[73]: mse3=mean_squared_error(y,re_y_pred)
mse3
```

```
[73]: 58823.49485131113
```

```
[74]: r2_score(y,re_y_pred)
```

```
[74]: 0.2186170394550624
```

0.0.23 Step19. [Select best model]. Tabulate MSE values of LR, KNNR and SGDR and select the model with the lowest MSE.

```
[76]: print("LR model ",mse_ln)
print("KNNR model ",mse)
print("SGDR model ",mse3)
```

```
LR model  96817.06978545067
KNNR model  70460.30507368421
SGDR model  58823.49485131113
```