

PML_Lab-5_205229118_Mahalakshmi.S

April 9, 2021

0.0.1 Lab5. Diabetes Classification using Logistic Regression

0.0.2 Step1. [Understand Data]. Using Pandas, import “diabetes.csv” file and print properties such as head, shape, columns, dtype, info and value_counts.

```
[2]: import pandas as pd
import csv
```

```
[3]: diab=pd.read_csv("diabetes.csv")
diab
```

```
[3]:      Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI   \
0                6      148             72             35         0  33.6
1                1       85             66             29         0  26.6
2                8      183             64              0         0  23.3
3                1       89             66             23        94  28.1
4                0      137             40             35       168  43.1
..            ...    ...             ...             ...    ...    ...
763             10      101             76             48       180  32.9
764              2      122             70             27         0  36.8
765              5      121             72             23       112  26.2
766              1      126             60              0         0  30.1
767              1       93             70             31         0  30.4
```

```
      DiabetesPedigreeFunction  Age  Outcome
0                0.627    50         1
1                0.351    31         0
2                0.672    32         1
3                0.167    21         0
4                2.288    33         1
..            ...    ...             ...
763             0.171    63         0
764             0.340    27         0
765             0.245    30         0
766             0.349    47         1
767             0.315    23         0
```

[768 rows x 9 columns]

```
[4]: diab.head()
```

```
[4]:   Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
0           6     148           72           35         0  33.6
1           1      85           66           29         0  26.6
2           8     183           64            0         0  23.3
3           1      89           66           23        94  28.1
4           0     137           40           35       168  43.1

      DiabetesPedigreeFunction  Age  Outcome
0                0.627     50         1
1                0.351     31         0
2                0.672     32         1
3                0.167     21         0
4                2.288     33         1
```

```
[5]: diab.shape
```

```
[5]: (768, 9)
```

```
[6]: df = pd.read_csv("diabetes.csv")
```

```
[7]: df
```

```
[7]:   Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
0           6     148           72           35         0  33.6
1           1      85           66           29         0  26.6
2           8     183           64            0         0  23.3
3           1      89           66           23        94  28.1
4           0     137           40           35       168  43.1
..          ...      ...             ...           ...      ...   ...
763         10     101           76           48       180  32.9
764          2     122           70           27         0  36.8
765          5     121           72           23       112  26.2
766          1     126           60            0         0  30.1
767          1      93           70           31         0  30.4

      DiabetesPedigreeFunction  Age  Outcome
0                0.627     50         1
1                0.351     31         0
2                0.672     32         1
3                0.167     21         0
4                2.288     33         1
..                  ...      ...       ...
763                0.171     63         0
764                0.340     27         0
765                0.245     30         0
```

```
766          0.349  47      1
767          0.315  23      0
```

```
[768 rows x 9 columns]
```

```
[8]: f = df.columns
      f
```

```
[8]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
          'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
          dtype='object')
```

```
[9]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies            768 non-null   int64
1   Glucose                768 non-null   int64
2   BloodPressure          768 non-null   int64
3   SkinThickness          768 non-null   int64
4   Insulin                768 non-null   int64
5   BMI                    768 non-null   float64
6   DiabetesPedigreeFunction 768 non-null   float64
7   Age                    768 non-null   int64
8   Outcome                768 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
[10]: df.dtypes
```

```
[10]: Pregnancies            int64
      Glucose                int64
      BloodPressure          int64
      SkinThickness          int64
      Insulin                int64
      BMI                    float64
      DiabetesPedigreeFunction float64
      Age                    int64
      Outcome                int64
      dtype: object
```

```
[11]: df.dtypes.value_counts()
```

```
[11]: int64      7
      float64    2
```

dtype: int64

0.0.3 Step2. [Build Logistic Regression Model]

Prepare X matrix (8 feature columns) and y vector (ie., Outcome column)

```
[12]: X=df.drop("Outcome", axis=1)
      y=df[["Outcome"]]
```

```
[13]: X
```

```
[13]:      Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
0              6      148             72             35         0  33.6
1              1       85             66             29         0  26.6
2              8      183             64              0         0  23.3
3              1       89             66             23        94  28.1
4              0      137             40             35       168  43.1
..          ...      ...             ...             ...      ...
763           10      101             76             48       180  32.9
764            2      122             70             27         0  36.8
765            5      121             72             23       112  26.2
766            1      126             60              0         0  30.1
767            1       93             70             31         0  30.4
```

```
      DiabetesPedigreeFunction  Age
0              0.627      50
1              0.351      31
2              0.672      32
3              0.167      21
4              2.288      33
..          ...      ...
763           0.171      63
764           0.340      27
765           0.245      30
766           0.349      47
767           0.315      23
```

[768 rows x 8 columns]

```
[14]: y
```

```
[14]:      Outcome
0              1
1              0
2              1
3              0
4              1
..          ...
```

```

763      0
764      0
765      0
766      1
767      0

```

```
[768 rows x 1 columns]
```

Split dataset with stratified shuffle split for training and testing as X_train, X_test, y_train, y_test (use 25% test size).

```
[15]: from sklearn.model_selection import StratifiedShuffleSplit
```

```
[16]: StratifiedShuffleSplit()
      shuf = StratifiedShuffleSplit(n_splits=4, test_size=0.25, random_state=0)
```

```
[17]: shuf.get_n_splits(X,y)
```

```
[17]: 4
```

Create LogisticRegression model, fit on training set and predict on test set

```
[18]: import warnings
      warnings.filterwarnings('ignore')
```

```
[19]: for train, test in shuf.split(X, y):
      X_train, X_test = X.iloc[train], X.iloc[test]
      y_train, y_test = y.iloc[train], y.iloc[test]
```

```
[20]: from sklearn.linear_model import LogisticRegression
```

```
[21]: logmodel = LogisticRegression()
```

```
[22]: logmodel.fit(X_train,y_train)
```

```
[22]: LogisticRegression()
```

```
[23]: y_predic = logmodel.predict(X_test)
      y_predic
```

```
[23]: array([0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0,
          0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1,
          1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0,
          0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0,
          1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1,
          0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0,
          1, 0, 1, 1, 1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0,
          0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0,
          0, 1, 1, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0], dtype=int64)
```

```
[24]: logmodel.score(X_train,y_train)
```

```
[24]: 0.7899305555555556
```

0.0.4 Step3. [Predict on a new sample]

Will this person become diabetic?. His details are given below.

```
new_person = [[6, 200, 90, 10, 25, 23.3, 0.672, 42]]
```

```
[25]: new_person = [[6, 200, 90, 10, 25, 23.3, 0.672, 42]]
```

```
[26]: print(logmodel.predict(new_person))
```

```
[1]
```

0.0.5 Step3. [Compute Classification Metrics]

Compute and print Accuracy, Precision, Recall and AUC scores

0.0.6 Precision

```
[27]: from sklearn.metrics import precision_score  
print(precision_score(y_test, y_predic))
```

```
0.6727272727272727
```

0.0.7 Recall

```
[28]: from sklearn.metrics import recall_score  
print(recall_score(y_test, y_predic))
```

```
0.5522388059701493
```

0.0.8 Accuracy

```
[29]: from sklearn.metrics import accuracy_score  
log_acscore=accuracy_score(y_test, y_predic)
```

```
[30]: log_acscore
```

```
[30]: 0.75
```

0.0.9 AUC scores

```
[31]: from sklearn.metrics import roc_auc_score
```

```
[33]: print(roc_auc_score(y_test, y_predic))
```

```
0.7041194029850747
```

0.0.10 Step4. [Understand Correlation]

Create confusion matrix between y_test and y_pred and plot confusion matrix values in a Heatmap. Explain the meaning of the 4 numbers you get.

```
[35]: from sklearn.metrics import confusion_matrix
      confu_matrix=confusion_matrix(y_test, y_predic)
      confu_matrix
```


```
[35]: array([[107,  18],
           [ 30,  37]], dtype=int64)
```

```
[37]: confu_accu_score = accuracy_score(y_test, y_predic)
      confu_accu_score
```

```
[37]: 0.75
```

```
[38]: import seaborn as sns
      sns.heatmap(confusion_matrix(y_test,y_predic) / len(y), cmap='YlGnBu',
      ↪annot=True)
```

```
[38]: <AxesSubplot:>
```



0.0.11 Step5. [Normalization using MinmaxScaler and rebuild LoR]

Now, normalize your X_train and X_test values using MinmaxScaler

```
[39]: from sklearn.preprocessing import MinMaxScaler
      scaler = MinMaxScaler()
```

```
X_trained_min = scaler.fit_transform(X_train)
X_tested_min = scaler.transform(X_test)
```

```
[40]: X_trained_min.shape
```

```
[40]: (576, 8)
```

```
[41]: X_tested_min.shape
```

```
[41]: (192, 8)
```

Create a new LogisticRegression model, fit on normalized training set and predict on the normalized test set

```
[42]: from sklearn.linear_model import LogisticRegression
logmodel1 = LogisticRegression()
```

```
[44]: logmodel1.fit(X_trained_min, y_train)
```

```
[44]: LogisticRegression()
```

```
[45]: y_predict = logmodel1.predict(X_tested_min)
y_predict
```

```
[45]: array([0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0,
          0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1,
          1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0,
          0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0,
          1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1,
          0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0,
          1, 0, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0,
          0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0,
          0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0], dtype=int64)
```

```
[46]: logmodel1.score(X_trained_min, y_train)
```

```
[46]: 0.7899305555555556
```

0.0.12 Compute and print Accuracy, Precision, Recall and AUC scores

0.0.13 Precision

```
[47]: from sklearn.metrics import precision_score
print(precision_score(y_test, y_predict))
```

```
0.673469387755102
```


0.0.14 Recall

```
[49]: from sklearn.metrics import recall_score  
print(recall_score(y_test, y_predict))
```

0.4925373134328358

0.0.15 Accuracy

```
[50]: from sklearn.metrics import accuracy_score  
min_accscore=accuracy_score(y_test, y_predict)
```

```
[51]: min_accscore
```

[51]: 0.7395833333333334

0.0.16 AUC scores

```
[52]: from sklearn.metrics import roc_auc_score
```

```
[56]: log_auc_sc=roc_auc_score(y_test, y_predict)  
log_auc1=('LoR minmax, AUC=',log_acscore)  
log_auc1
```

[56]: ('LoR minmax, AUC=', 0.75)

0.0.17 Step6. [Normalization using StandardScaler and rebuild LoR]

Repeat Step5 with StandardScaler

```
[57]: from sklearn.preprocessing import StandardScaler  
scaler = StandardScaler()  
X_trained_stand = scaler.fit_transform(X_train)  
X_tested_stand = scaler.transform(X_test)
```

```
[59]: X_trained_stand.shape
```

[59]: (576, 8)

```
[60]: X_tested_stand.shape
```

[60]: (192, 8)

```
[61]: from sklearn.linear_model import LogisticRegression  
logmodel2 = LogisticRegression()
```

```
[63]: logmodel2.fit(X_trained_stand, y_train)
```

[63]: LogisticRegression()

```
[66]: y_predict_stand = logmodel2.predict(X_tested_stand)
      y_predict_stand
```

```
[66]: array([0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0,
          0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1,
          1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0,
          0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0,
          1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1,
          0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0,
          1, 0, 1, 1, 1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 0, 0,
          0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0,
          0, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0])
```

```
[67]: logmodel2.score(X_trained_stand, y_train)
```

```
[67]: 0.7795138888888888
```

0.0.18 Precision score

```
[68]: from sklearn.metrics import precision_score
      print(precision_score(y_test, y_predict_stand))
```

```
0.6851851851851852
```

0.0.19 Recall score

```
[69]: from sklearn.metrics import recall_score
      print(recall_score(y_test, y_predict_stand))
```

```
0.5522388059701493
```

0.0.20 Accuracy

```
[70]: from sklearn.metrics import accuracy_score
      stand_accscore=accuracy_score(y_test, y_predict_stand)
```

```
[71]: stand_accscore
```

```
[71]: 0.7552083333333334
```

0.0.21 AUC score

```
[72]: from sklearn.metrics import roc_auc_score
```

```
[73]: stand_aucscore=roc_auc_score(y_test, y_predict_stand)
      stand_auc3=('AUC=',stand_aucscore)
      stand_auc3
```

```
[73]: ('AUC=', 0.7081194029850746)
```

0.0.22 Among the 3 models, which model gives better classification scores?

```
[74]: print('Logistic Regression Model:',log_acscore)
      print('MinmaxScaler:',min_accscore)
      print('StandardScaler:',stand_accscore)
```

```
Logistic Regression Model: 0.75
MinmaxScaler: 0.7395833333333334
StandardScaler: 0.7552083333333334
```

0.0.23 Step7. [Plot ROC curve]

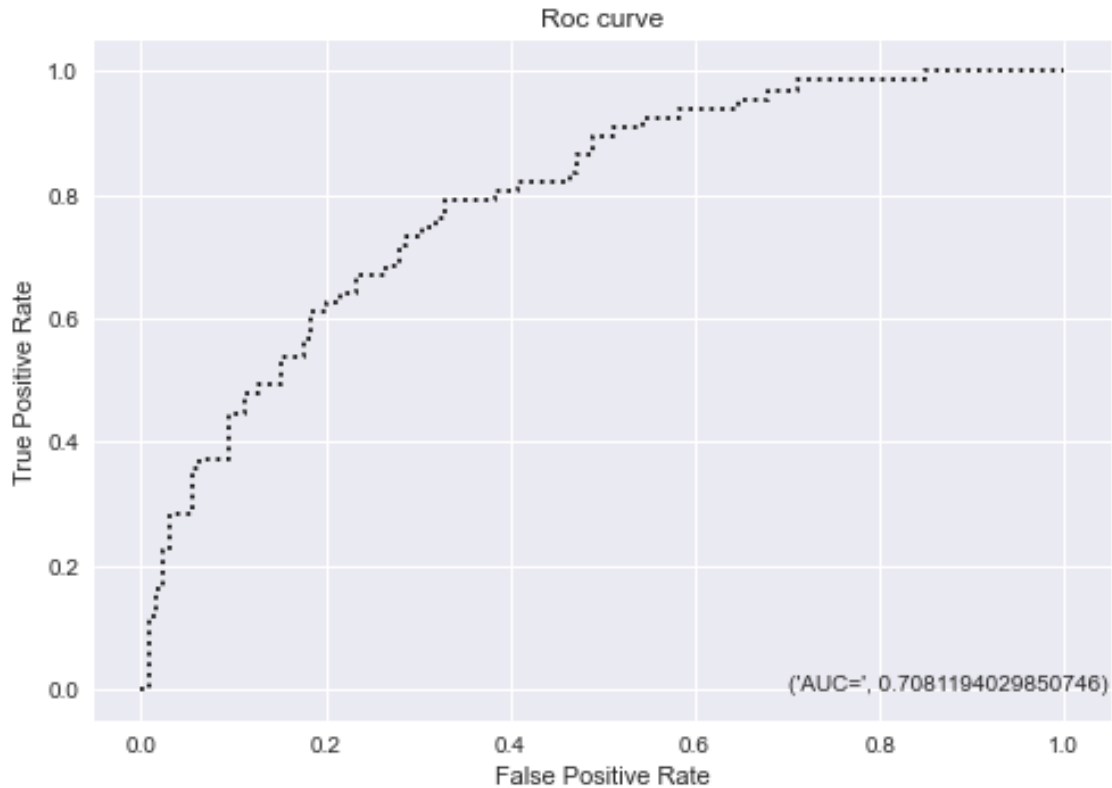
Plot ROC curve as shown below. You can use the MinmaxScaler scaled values of `X_test` for computing `predict_proba()` score.¶

```
[75]: from sklearn.metrics import roc_curve
```

```
[77]: pred_prb3 = logmodel1.predict_proba(X_tested_min)
      fprb3, tprb3, threshold3 = roc_curve(y_test,pred_prb3[:,1], pos_label=1)
```

```
[78]: import matplotlib.pyplot as plt
      plt.style.use('seaborn')
      plt.annotate(xy=[0.7,0], s=stand_auc3)
      plt.plot(fpr1, tpr1, linestyle=':', color='black',label='Logistic Regression')
      plt.title('Roc curve')
      plt.xlabel('False Positive Rate')
      plt.ylabel('True Positive Rate')
```

```
[78]: Text(0, 0.5, 'True Positive Rate')
```



0.0.24 Step8. [Comparison with KNN classifier].

Create a KNN classifier with default values, fit on the scaled X using MinmaxScaler, predict and print classification metric scores.

```
[79]: from sklearn.neighbors import KNeighborsClassifier
```

```
[80]: logmodel3 = KNeighborsClassifier(n_neighbors=3)
```

```
[81]: logmodel3.fit(X_trained_min,y_train)
```

```
[81]: KNeighborsClassifier(n_neighbors=3)
```

```
[82]: knn_y_pred = logmodel3.predict(X_tested_min)
      knn_y_pred
```

```
[82]: array([1, 0, 1, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0,
            0, 0, 1, 0, 1, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1,
            1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0,
            0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 1, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0,
            1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 1, 0, 0,
            0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0,
            0, 0, 1, 1, 1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0,
```

```
0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 0,
0, 1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1], dtype=int64)
```

0.0.25 Precision score

```
[83]: from sklearn.metrics import precision_score
      print(precision_score(y_test, knn_y_pred))
```

```
0.6271186440677966
```

0.0.26 Recall score

```
[84]: from sklearn.metrics import recall_score
      print(recall_score(y_test, knn_y_pred))
```

```
0.5522388059701493
```

0.0.27 Accuracy

```
[85]: from sklearn.metrics import accuracy_score
      knn_accscore=accuracy_score(y_test, knn_y_pred)
      knn_accscore
```

```
[85]: 0.7291666666666666
```

0.0.28 AUC score

```
[87]: from sklearn.metrics import roc_auc_score
      knn_aucscore=roc_auc_score(y_test, knn_y_pred)
      knn_auc2=('KNN minmax, AUC=',knn_aucscore)
      knn_auc2
```

```
[87]: ('KNN minmax, AUC=', 0.6881194029850747)
```

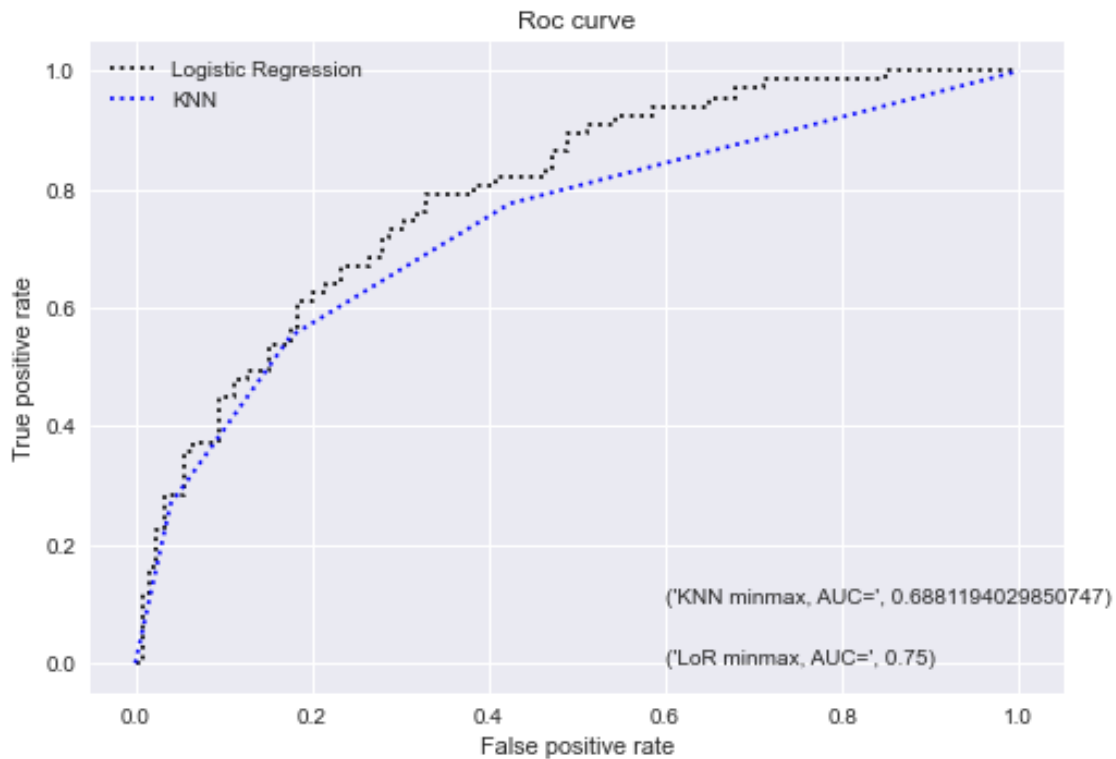
0.0.29 Step9. [Update ROC curve]

Update your ROC curve, this time, with one more curve of KNN classifier

```
[89]: pred_prb2 = logmodel3.predict_proba(X_tested_min)
      fprb3,tprb3,threshold1 = roc_curve(y_test, pred_prb3[:,1],pos_label=1)
      fprb2,tprb2,threshold2 = roc_curve(y_test, pred_prb2[:,1],pos_label=1)

[91]: plt.plot(fprb3,tprb3,linestyle=':',color='black',label='Logistic Regression')
      plt.plot(fprb2,tprb2,linestyle=':',color='blue',label='KNN')
      plt.annotate(xy=[0.6,0.1], s=knn_auc2)
      plt.annotate(xy=[0.6,0], s=log_auc1)
      plt.legend(loc='best')
      plt.title('Roc curve')
      plt.xlabel('False positive rate')
      plt.ylabel('True positive rate')
```

```
[91]: Text(0, 0.5, 'True positive rate')
```



0.0.30 Step10. [Regularization]

In order to reduce overfitting of your data, you will use `LogisticRegressionCV` model with L1 and L2 regularization parameters. Create both models using the following statements

```
[92]: from sklearn.linear_model import LogisticRegressionCV
```

```
[93]: logmodel4 = LogisticRegressionCV(Cs=10, cv=4, penalty='l1', solver='liblinear')
logmodel5 = LogisticRegressionCV(Cs=10, cv=4, penalty='l2')
```

Perform fit using `MinmaxScaler` scaled values and predict¶

```
[94]: print(logmodel4.fit(X_trained_min,y_train))
print(logmodel5.fit(X_trained_min,y_train))
```

```
LogisticRegressionCV(cv=4, penalty='l1', solver='liblinear')
LogisticRegressionCV(cv=4)
```

```
[97]: logrg_y_pred5 = logmodel4.predict(X_tested_min)
logrg_y_pred6 = logmodel5.predict(X_tested_min)
print('Logistic RegressionCV L1:\n',logrg_y_pred5)
print('#####')
```

```
print('Logistic RegressionCV L2:\n',logrg_y_pred6)
```

Logistic RegressionCV L1:

```
[0 0 1 1 0 1 1 0 0 1 0 0 0 1 0 0 0 1 0 0 0 0 0 0 0 1 0 1 0 0 1 0 0 0 0 0 0
0 0 0 0 1 0 1 1 0 1 0 1 0 0 0 1 0 0 1 0 0 0 1 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0
1 0 0 0 0 1 1 0 0 0 0 0 1 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 1 0 0 1 0
0 0 1 1 0 0 0 0 0 1 0 0 1 1 0 0 0 0 0 0 0 0 1 0 1 1 1 1 0 0 0 0 0 0 1 0 0 0
1 1 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 1 0 1 0 0 1 0 0 0 0 0 1 0 0 1 1 0 0
0 0 0 1 0 0 0]
```

```
#####
#####
```

Logistic RegressionCV L2:

```
[0 0 1 1 0 1 1 0 0 1 0 0 1 1 0 0 0 1 0 0 0 0 0 0 0 1 0 1 0 0 1 0 0 0 0 0 0
0 0 0 0 1 0 1 1 0 1 0 1 0 0 0 1 0 0 1 0 0 0 1 0 0 1 0 0 0 0 0 0 0 0 0 0 0 1
1 0 0 0 0 1 1 0 0 0 0 0 1 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 1 0 0 1 0
0 0 1 1 0 0 0 0 0 1 0 0 1 1 0 0 0 0 0 0 0 0 1 0 1 1 1 1 0 0 0 1 0 0 1 0 0 0
1 1 0 0 0 0 0 1 0 0 0 0 0 0 0 1 0 0 0 1 0 1 0 0 1 0 0 0 0 0 1 0 0 1 1 0 0
0 0 0 1 0 0 0]
```

```
[98]: from sklearn.metrics import roc_auc_score
logrgpred_auc=roc_auc_score(y_test,logrg_y_pred5)
logrgpred4_auc=('LoR L1 minmax, AUC=',logrgpred_auc)
logrgpred4_auc
```

```
[98]: ('LoR L1 minmax, AUC=', 0.6748059701492537)
```

```
[99]: from sklearn.metrics import roc_auc_score
logrgpred_auc1=roc_auc_score(y_test, logrg_y_pred6)
logrgpred5_auc=('LoR L2 minmax, AUC=',logrgpred_auc1)
logrgpred5_auc
```

```
[99]: ('LoR L2 minmax, AUC=', 0.6931940298507463)
```

0.0.31 Step11. [Update ROC curve]

Update your ROC curve, this time, with two more curves

```
[101]: pred_prb7 = logmodel4.predict_proba(X_tested_min)
pred_prb8 = logmodel5.predict_proba(X_tested_min)
fprb2,tprb2,threshold2 = roc_curve(y_test, pred_prb2[:,1],pos_label=1)
fprb3,tprb3,threshold1 = roc_curve(y_test, pred_prb3[:,1],pos_label=1)
fprb5,tprb5,thresh3 = roc_curve(y_test, pred_prb7[:,1],pos_label=1)
fprb6,tprb6,thresh4 = roc_curve(y_test, pred_prb8[:,1],pos_label=1)

[102]: plt.plot(fprb2,tprb2,linestyle='-',color='midnightblue', label='Logistic_
→Regression')
plt.plot(fprb3,tprb3,linestyle='-',color='black', label='KNN')
plt.plot(fprb5,tprb5,linestyle='-',color='red', label='LoR L1')
plt.plot(fprb6,tprb6,linestyle='-',color='forestgreen', label='LoR L2')
```

```

plt.annotate(xy=[0.49,0.3], s=knn_auc2)
plt.annotate(xy=[0.49,0.2], s=log_auc1)
plt.annotate(xy=[0.49,0.1], s=logrgpred4_auc)
plt.annotate(xy=[0.49,0], s=logrgpred5_auc)
plt.legend(loc='best')
plt.title('ROC curve')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')

```

[102]: Text(0, 0.5, 'True Positive Rate')

