

PML_lab-7_205229118_Mahalakshmi.S

April 27, 2021

0.1 Lab7. Loan Approval Classification using SVM

Objectives In this lab, you will build a classification model to classify the loan applicants into eligible applicants or not eligible applicants using Support Vector Machine.

Step1. [Understand Data]. Using Pandas, import “train_loan.csv” file and print properties such as head, shape, columns, dtype, info and value_counts

```
[1]: import pandas as pd
import matplotlib.pyplot as plt

tralon_data = pd.read_csv('train_loan.csv')
tralon_data.head()
```

```
[1]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
0	LP001002	Male	No	0	Graduate	No	
1	LP001003	Male	Yes	1	Graduate	No	
2	LP001005	Male	Yes	0	Graduate	Yes	
3	LP001006	Male	Yes	0	Not Graduate	No	
4	LP001008	Male	No	0	Graduate	No	

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	\
0	5849	0.0	NaN	360.0	
1	4583	1508.0	128.0	360.0	
2	3000	0.0	66.0	360.0	
3	2583	2358.0	120.0	360.0	
4	6000	0.0	141.0	360.0	

	Credit_History	Property_Area	Loan_Status
0	1.0	Urban	Y
1	1.0	Rural	N
2	1.0	Urban	Y
3	1.0	Urban	Y
4	1.0	Urban	Y

```
[2]: tralon_data.shape
```

```
[2]: (614, 13)
```

```
[3]: tralon_data.columns
```

```
[3]: Index(['Loan_ID', 'Gender', 'Married', 'Dependents', 'Education',  
        'Self_Employed', 'ApplicantIncome', 'CoapplicantIncome', 'LoanAmount',  
        'Loan_Amount_Term', 'Credit_History', 'Property_Area', 'Loan_Status'],  
        dtype='object')
```

```
[4]: tralon_data.dtypes
```

```
[4]: Loan_ID           object  
     Gender           object  
     Married          object  
     Dependents       object  
     Education        object  
     Self_Employed    object  
     ApplicantIncome  int64  
     CoapplicantIncome float64  
     LoanAmount       float64  
     Loan_Amount_Term float64  
     Credit_History   float64  
     Property_Area     object  
     Loan_Status      object  
     dtype: object
```

```
[5]: tralon_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 614 entries, 0 to 613  
Data columns (total 13 columns):  
#   Column                Non-Null Count  Dtype  
---  ---  
0   Loan_ID               614 non-null   object  
1   Gender                601 non-null   object  
2   Married               611 non-null   object  
3   Dependents            599 non-null   object  
4   Education             614 non-null   object  
5   Self_Employed         582 non-null   object  
6   ApplicantIncome       614 non-null   int64  
7   CoapplicantIncome     614 non-null   float64  
8   LoanAmount            592 non-null   float64  
9   Loan_Amount_Term      600 non-null   float64  
10  Credit_History        564 non-null   float64  
11  Property_Area         614 non-null   object  
12  Loan_Status           614 non-null   object  
dtypes: float64(4), int64(1), object(8)  
memory usage: 62.5+ KB
```

```
[6]: tralon_data.value_counts
```

```
[6]: <bound method DataFrame.value_counts of
      Education Self_Employed \
      Loan_ID Gender Married Dependents
```

0	LP001002	Male	No	0	Graduate	No
1	LP001003	Male	Yes	1	Graduate	No
2	LP001005	Male	Yes	0	Graduate	Yes
3	LP001006	Male	Yes	0	Not Graduate	No
4	LP001008	Male	No	0	Graduate	No
..
609	LP002978	Female	No	0	Graduate	No
610	LP002979	Male	Yes	3+	Graduate	No
611	LP002983	Male	Yes	1	Graduate	No
612	LP002984	Male	Yes	2	Graduate	No
613	LP002990	Female	No	0	Graduate	Yes

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	\
0	5849	0.0	NaN	360.0	
1	4583	1508.0	128.0	360.0	
2	3000	0.0	66.0	360.0	
3	2583	2358.0	120.0	360.0	
4	6000	0.0	141.0	360.0	
..	
609	2900	0.0	71.0	360.0	
610	4106	0.0	40.0	180.0	
611	8072	240.0	253.0	360.0	
612	7583	0.0	187.0	360.0	
613	4583	0.0	133.0	360.0	

	Credit_History	Property_Area	Loan_Status
0	1.0	Urban	Y
1	1.0	Rural	N
2	1.0	Urban	Y
3	1.0	Urban	Y
4	1.0	Urban	Y
..
609	1.0	Rural	Y
610	1.0	Rural	Y
611	1.0	Urban	Y
612	1.0	Urban	Y
613	0.0	Semiurban	N

```
[614 rows x 13 columns]>
```

0.1.1 Step2. [Data Cleaning]

0.1.2 Replace numbers as string by integer in “Dependents” column

```
[7]: tralon_data.Dependents.value_counts()
```

```
[7]: 0      345
      1      102
      2      101
      3+      51
      Name: Dependents, dtype: int64
```

```
[8]: #Replace numbers as string by integer in 'Dependents' column
def string(x):
    if x == '0':
        return 'bad'
    elif x == '1':
        return 'average'
    elif x == '2':
        return 'good'
    else:
        return 'excellent'
```

```
[9]: tralon_data['Dependents'] = tralon_data['Dependents'].apply(string)
```

0.1.3 Fill missing data in categorical columns (Gender, Married, Dependents, Education, Self_Employed, Credit_History) by its mode value

0.1.4 Handle missing values in numerical columns

```
[10]: tralon_data.isna().sum()
```

```
[10]: Loan_ID      0
      Gender      13
      Married      3
      Dependents  0
      Education   0
      Self_Employed  32
      ApplicantIncome  0
      CoapplicantIncome  0
      LoanAmount    22
      Loan_Amount_Term  14
      Credit_History  50
      Property_Area  0
      Loan_Status    0
      dtype: int64
```

```
[11]: #categorical
tralon_data['Gender'].fillna(tralon_data['Gender'].mode()[0], inplace=True)
tralon_data['Married'].fillna(tralon_data['Married'].mode()[0], inplace=True)
tralon_data['Dependents'].fillna(tralon_data['Dependents'].mode()[0],
    ↳inplace=True)
tralon_data['Loan_Amount_Term'].fillna(tralon_data['Loan_Amount_Term'].
    ↳mode()[0], inplace=True)
tralon_data['Credit_History'].fillna(tralon_data['Credit_History'].mode()[0],
    ↳inplace=True)
tralon_data['Self_Employed'].fillna(tralon_data['Self_Employed'].mode()[0],
    ↳inplace=True)
#numerical
tralon_data['LoanAmount'].fillna(tralon_data['LoanAmount'].mean(), inplace=True)
```

0.1.5 Drop Loan_ID column

```
[12]: tralon_data=tralon_data.drop(['Loan_ID'],axis=1)
```

0.1.6 Step3. [OPTIONAL: Exploratory Data Analysis - Who got their loan approved]

Draw count plot for

Married?

Dependants?

Graduates?

Self-employed?


```
[13]: plt.subplot(231)
tralon_data['Married'].value_counts().plot(kind='bar',title='Gender',figsize =(
    ↳(20,10))

plt.subplot(232)
tralon_data['Dependents'].value_counts().plot(kind='bar',title='Dependents')

plt.subplot(233)
tralon_data['Education'].value_counts().plot(kind='bar',title='Education')

plt.subplot(234)
tralon_data['Self_Employed'].value_counts().
    ↳plot(kind='bar',title='Self_Employed')

plt.show()
```



0.1.7 Step4. [Extract X and y] from the dataframe

```
[14]: X = tralon_data.drop(['Loan_Status'],axis=1)
      y = tralon_data.Loan_Status
```

0.1.8 Step5. [One Hot Encoding]

Perform OHE on categorical columns, use this method: `X = pd.get_dummies(X)`

```
[15]: import warnings
      warnings.filterwarnings('ignore')
```

```
[16]: X = pd.get_dummies(X)
```

0.1.9 Step6. [Model Building]

Split X and y for training and testing

```
[17]: from sklearn.model_selection import train_test_split
      X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.
      ↳3,random_state=0)
```

Using StandardScaler, fit_transform on X_train and transform on X_test values

```
[18]: from sklearn.preprocessing import StandardScaler
      st_sc = StandardScaler()
```

```
[22]: stand_sc = st_sc.fit_transform(X_train)
      stand_sc1 = st_sc.transform(X_test)
```

create LinearSVC model, train and test

```
[23]: from sklearn.svm import LinearSVC
lin_model = LinearSVC()
lin_model.fit(stand_sc,y_train)
```

```
[23]: LinearSVC()
```

```
[25]: linsvc_y_predict = lin_model.predict(stand_sc1)
linsvc_y_predict
```

```
[25]: array(['Y', 'Y', 'Y', 'Y', 'Y', 'N', 'Y', 'Y', 'N', 'Y', 'Y', 'Y', 'Y',
          'Y', 'Y', 'Y', 'Y', 'Y', 'N', 'N', 'Y', 'Y', 'Y', 'Y', 'Y',
          'Y', 'Y', 'N', 'N', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'N', 'Y', 'Y', 'Y',
          'Y', 'Y', 'N', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y',
          'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y',
          'Y', 'N', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y',
          'Y', 'Y', 'N', 'Y', 'N', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y',
          'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'N', 'N', 'Y',
          'Y', 'Y', 'Y', 'Y', 'N', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y',
          'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y',
          'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y',
          'Y', 'N', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y', 'Y',
          'Y', 'Y', 'Y'], dtype=object)
```

print accuracy value

```
[26]: from sklearn.metrics import accuracy_score,confusion_matrix,
      ↪classification_report
accuracy_score(y_test,linsvc_y_predict)
```

```
[26]: 0.8324324324324325
```

Print confusion matrix between y_test and y_pred

```
[27]: confusion_matrix(y_test,linsvc_y_predict)
```

```
[27]: array([[ 22,  29],
          [  2, 132]], dtype=int64)
```

Print classification_report

```
[28]: print(classification_report(y_test,linsvc_y_predict))
```

	precision	recall	f1-score	support
N	0.92	0.43	0.59	51
Y	0.82	0.99	0.89	134
accuracy			0.83	185

macro avg	0.87	0.71	0.74	185
weighted avg	0.85	0.83	0.81	185

0.1.10 Step7. [Performance Comparisons]

1. Compare the performance of LinearSVC against LogisticRegression

```
[29]: from sklearn.linear_model import LogisticRegression

logreg= LogisticRegression()
logreg.fit(stand_sc,y_train)
logreg_y_predict = logreg.predict(stand_sc1)

from sklearn.svm import LinearSVC

ln_svc = LinearSVC()
ln_svc.fit(stand_sc,y_train)
lnsvc_y_predict = ln_svc.predict(stand_sc1)

print("LogisticRegression:",accuracy_score(y_test,logreg_y_predict))
print("LinearSVC           :",accuracy_score(y_test,lnsvc_y_predict))
```

LogisticRegression: 0.8324324324324325


LinearSVC : 0.8324324324324325

```
[30]: from sklearn import svm,model_selection
models = []
models.append(('SVC', LinearSVC()))
models.append(('LR', LogisticRegression()))

# evaluate each model in turn
results = []
names=[]
scoring = 'accuracy'
for name,model in models:
    kfold = model_selection.KFold(n_splits=10)
    cv_results = model_selection.cross_val_score(model, X_train, y_train,
    ↪cv=kfold, scoring=scoring)
    results.append(cv_results)
    names.append(name)
    msg = "%s: %f (%f)" % (name, cv_results.mean(), cv_results.std())
    print(msg)
# boxplot algorithm comparison
fig = plt.figure()
fig.suptitle('Comparison between different MLAs')
ax = fig.add_subplot(111)
plt.boxplot(results)
plt.show()
```


SVC: 0.580066 (0.155858)
LR: 0.797231 (0.070487)

Comparison between different MLAs



2. Compare the performance of LinearSVC against SGDClassifier

```
[31]: from sklearn.linear_model import SGDClassifier

modelss = []
modelss.append(('SVC', LinearSVC()))
modelss.append(('SGD', SGDClassifier()))

# evaluate each model in turn
results = []
names=[]
scoring = 'accuracy'
for name,model in modelss:
    kfold = model_selection.KFold(n_splits=10)
    cv_results = model_selection.cross_val_score(model, X_train, y_train,
    ↪cv=kfold, scoring=scoring)
    results.append(cv_results)
    names.append(name)
    msg = "%s: %f (%f)" % (name, cv_results.mean(), cv_results.std())
    print(msg)
# boxplot algorithm comparison
```


```

fig = plt.figure()
fig.suptitle('Comparison between different MLAs')
ax = fig.add_subplot(111)
plt.boxplot(results)
plt.show()

```

SVC: 0.543189 (0.160649)

SGD: 0.585050 (0.117652)



```

[40]: from sklearn.linear_model import SGDClassifier

sgdc = SGDClassifier()
sgdc.fit(stand_sc,y_train)
sgdc1_y_predict = sgdc.predict(stand_sc1)

from sklearn.svm import LinearSVC

ln_svc = LinearSVC()
ln_svc.fit(stand_sc,y_train)
lnsvc_y_predict = ln_svc.predict(stand_sc1)

print("SGDClassifier:", accuracy_score(y_test,sgdc1_y_predict))
print("LinearSVC      :",accuracy_score(y_test,lnsvc_y_predict))

```

```
SGDClassifier: 0.7675675675675676
LinearSVC      : 0.8324324324324325
```

3. Compare LinearSVC against SVC with various kernels such as 'linear', 'poly', 'rbf' and 'sigmoid'

```
[41]: from sklearn.svm import SVC

ln_svc = LinearSVC()
ln_svc.fit(stand_sc,y_train)
lnsvc_y_predict = ln_svc.predict(stand_sc1)

poly_svc = svm.SVC(kernel='poly', C = 1.0)
poly_svc.fit(stand_sc,y_train)
polsvc_y_predict=poly_svc.predict(stand_sc1)

rbf_svc = svm.SVC(kernel='rbf', C = 1.0)
rbf_svc.fit(stand_sc,y_train)
rbfsvc_y_predict=rbf_svc.predict(stand_sc1)

sig_svc = svm.SVC(kernel='sigmoid', C = 1.0)
sig_svc.fit(stand_sc,y_train)
sigsvc_y_predict=sig_svc.predict(stand_sc1)

print("LinearSVC      :",accuracy_score(y_test,lnsvc_y_predict))
print("poly SVC       :",accuracy_score(y_test,polsvc_y_predict))
print("rbf SVC        :",accuracy_score(y_test,rbfsvc_y_predict))
print("Sigmoid SVC    :",accuracy_score(y_test,sigsvc_y_predict))

LinearSVC      : 0.8324324324324325
poly SVC       : 0.8162162162162162
rbf SVC        : 0.8324324324324325
Sigmoid SVC    : 0.8054054054054054
```

```
[42]: models1 = []
models1.append(('SVC', LinearSVC()))
models1.append(('SVC POLY', svm.SVC(kernel='poly', C = 1.0)))
models1.append(('SVC rbf', svm.SVC(kernel='rbf', C = 1.0)))
models1.append(('SVC POLY', svm.SVC(kernel='sigmoid', C = 1.0)))


# evaluate each model in turn
results = []
names=[]
scoring = 'accuracy'
for name,model in models1:
```

```

kfold = model_selection.KFold(n_splits=10)
cv_results = model_selection.cross_val_score(model, X_train, y_train,
cv=kfold, scoring=scoring)
results.append(cv_results)
names.append(name)
msg = "%s: %f (%f)" % (name, cv_results.mean(), cv_results.std())
print(msg)
# boxplot algorithm comparison
fig = plt.figure()
fig.suptitle('Comparison between different MLAs')
ax = fig.add_subplot(111)
plt.boxplot(results)
plt.show()

```

SVC: 0.506091 (0.172120)
SVC POLY: 0.671096 (0.047891)
SVC rbf: 0.666445 (0.055735)
SVC POLY: 0.603710 (0.092483)



4. Interpret the results

```

[44]: import numpy as np
from sklearn.metrics import roc_curve
from sklearn.metrics import precision_score

```

```

from sklearn.metrics import recall_score
from sklearn.metrics import auc

ML = [model,logreg,sgdc,poly_svc,rbf_svc,sig_svc]
ML_columns = []
ML_compare = pd.DataFrame(columns = ML_columns)

row_index = 0
for alg in ML:

    predicted = alg.fit(stand_sc, y_train).predict(stand_sc1)
    predicted=np.where(predicted=='Y',1,0)
    y_testb=np.where(y_test=='Y',1,0)
    fp, tp, th = roc_curve(y_testb, predicted)
    ML_name = alg.__class__.__name__
    ML_compare.loc[row_index, 'ML used'] = ML_name
    ML_compare.loc[row_index, 'Train Accuracy'] = round(alg.
→score(stand_sc,y_train), 4)
    ML_compare.loc[row_index, 'Test Accuracy'] = round(alg.
→score(stand_sc1,y_test), 4)
    ML_compare.loc[row_index, 'Precission'] = precision_score(y_testb,
→predicted)
    ML_compare.loc[row_index, 'Recall'] = recall_score(y_testb, predicted)
    ML_compare.loc[row_index, 'AUC'] = auc(fp, tp)
    row_index+=1

ML_compare

```

```

[44]:
      ML used  Train Accuracy  Test Accuracy  Precission  Recall \
0          SVC           0.7506           0.8054    0.810127  0.955224
1  LogisticRegression           0.8042           0.8324    0.819876  0.985075
2    SGDClassifier           0.7669           0.7892    0.818792  0.910448
3          SVC           0.8368           0.8162    0.820513  0.955224
4          SVC           0.8135           0.8324    0.819876  0.985075
5          SVC           0.7506           0.8054    0.810127  0.955224

      AUC
0  0.683494
1  0.708224
2  0.690518
3  0.703102
4  0.708224
5  0.683494

```

```

[45]: import seaborn as sns
      # Creating plot to show the train accuracy
      plt.subplots(figsize=(8,4))

```

```

sns.barplot(x="ML used", y="Train_Accuracy", data=ML_compare, palette='hot', edgecolor=sns.
↳color_palette('dark',7))
plt.xticks(rotation=90)
plt.title('ML Train Accuracy Comparison')
plt.show()

```

