

# PML\_Lab-9\_205229118\_Mahalakshmi.S

May 27, 2021

## 0.0.1 Lab9. Employee Hopping Prediction using Random Forests

```
[16]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import tree
from sklearn.metrics import precision_score, \
    recall_score, accuracy_score, roc_auc_score, classification_report, f1_score
from sklearn.tree import export_graphviz
from sklearn.preprocessing import LabelEncoder
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
import warnings
warnings.filterwarnings('ignore')
```

## 0.0.2 Step1. [Understand Data].

Using Pandas, import “Employee\_Hopping.csv” file and print properties such as head, shape, columns, dtype, info and value\_counts.

```
[17]: emp = pd.read_csv("Employee_Hopping.csv")
```

```
[18]: emp.head()
```

```
[18]:
```

	Age	Attrition	BusinessTravel	DailyRate	Department	\
0	41	Yes	Travel_Rarely	1102		Sales
1	49	No	Travel_Frequently	279	Research & Development	
2	37	Yes	Travel_Rarely	1373	Research & Development	
3	33	No	Travel_Frequently	1392	Research & Development	
4	27	No	Travel_Rarely	591	Research & Development	

	DistanceFromHome	Education	EducationField	EmployeeCount	EmployeeNumber	\
0		1	2 Life Sciences	1		1
1		8	1 Life Sciences	1		2
2		2	2 Other	1		4
3		3	4 Life Sciences	1		5
4		2	1 Medical	1		7

	...	RelationshipSatisfaction	StandardHours	StockOptionLevel	\
0	...	1	80	0	
1	...	4	80	1	
2	...	2	80	0	
3	...	3	80	0	
4	...	4	80	1	

	TotalWorkingYears	TrainingTimesLastYear	WorkLifeBalance	YearsAtCompany	\
0	8	0	1	6	
1	10	3	3	10	
2	7	3	3	0	
3	8	3	3	8	
4	6	3	3	2	

	YearsInCurrentRole	YearsSinceLastPromotion	YearsWithCurrManager
0	4	0	5
1	7	1	7
2	0	0	0
3	7	3	0
4	2	2	2

[5 rows x 35 columns]

```
[19]: emp.shape
```

```
[19]: (1470, 35)
```

```
[20]: emp.columns
```

```
[20]: Index(['Age', 'Attrition', 'BusinessTravel', 'DailyRate', 'Department',
        'DistanceFromHome', 'Education', 'EducationField', 'EmployeeCount',
        'EmployeeNumber', 'EnvironmentSatisfaction', 'Gender', 'HourlyRate',
        'JobInvolvement', 'JobLevel', 'JobRole', 'JobSatisfaction',
        'MaritalStatus', 'MonthlyIncome', 'MonthlyRate', 'NumCompaniesWorked',
        'Over18', 'OverTime', 'PercentSalaryHike', 'PerformanceRating',
        'RelationshipSatisfaction', 'StandardHours', 'StockOptionLevel',
        'TotalWorkingYears', 'TrainingTimesLastYear', 'WorkLifeBalance',
        'YearsAtCompany', 'YearsInCurrentRole', 'YearsSinceLastPromotion',
        'YearsWithCurrManager'],
        dtype='object')
```

```
[21]: emp.dtypes
```

```
[21]: Age                int64
Attrition              object
BusinessTravel         object
DailyRate              int64
```

Department	object
DistanceFromHome	int64
Education	int64
EducationField	object
EmployeeCount	int64
EmployeeNumber	int64
EnvironmentSatisfaction	int64
Gender	object
HourlyRate	int64
JobInvolvement	int64
JobLevel	int64
JobRole	object
JobSatisfaction	int64
MaritalStatus	object
MonthlyIncome	int64
MonthlyRate	int64
NumCompaniesWorked	int64
Over18	object
OverTime	object
PercentSalaryHike	int64
PerformanceRating	int64
RelationshipSatisfaction	int64
StandardHours	int64
StockOptionLevel	int64
TotalWorkingYears	int64
TrainingTimesLastYear	int64
WorkLifeBalance	int64
YearsAtCompany	int64
YearsInCurrentRole	int64
YearsSinceLastPromotion	int64
YearsWithCurrManager	int64
dtype:	object

```
[22]: emp.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 35 columns):
#   Column              Non-Null Count  Dtype
---  -
0   Age                 1470 non-null   int64
1   Attrition           1470 non-null   object
2   BusinessTravel      1470 non-null   object
3   DailyRate           1470 non-null   int64
4   Department          1470 non-null   object
5   DistanceFromHome    1470 non-null   int64
6   Education            1470 non-null   int64
7   EducationField       1470 non-null   object
```

```

8   EmployeeCount      1470 non-null   int64
9   EmployeeNumber     1470 non-null   int64
10  EnvironmentSatisfaction  1470 non-null   int64
11  Gender              1470 non-null   object
12  HourlyRate          1470 non-null   int64
13  JobInvolvement      1470 non-null   int64
14  JobLevel            1470 non-null   int64
15  JobRole             1470 non-null   object
16  JobSatisfaction     1470 non-null   int64
17  MaritalStatus       1470 non-null   object
18  MonthlyIncome       1470 non-null   int64
19  MonthlyRate         1470 non-null   int64
20  NumCompaniesWorked  1470 non-null   int64
21  Over18              1470 non-null   object
22  OverTime            1470 non-null   object
23  PercentSalaryHike   1470 non-null   int64
24  PerformanceRating   1470 non-null   int64
25  RelationshipSatisfaction 1470 non-null   int64
26  StandardHours       1470 non-null   int64
27  StockOptionLevel    1470 non-null   int64
28  TotalWorkingYears   1470 non-null   int64
29  TrainingTimesLastYear 1470 non-null   int64
30  WorkLifeBalance     1470 non-null   int64
31  YearsAtCompany      1470 non-null   int64
32  YearsInCurrentRole  1470 non-null   int64
33  YearsSinceLastPromotion 1470 non-null   int64
34  YearsWithCurrManager 1470 non-null   int64

```

dtypes: int64(26), object(9)

memory usage: 402.1+ KB

```
[23]: emp.value_counts
```

```

[23]: <bound method DataFrame.value_counts of
DailyRate      Department \
0      41      Yes      Travel_Rarely      1102      Sales
1      49      No      Travel_Frequently      279      Research & Development
2      37      Yes      Travel_Rarely      1373      Research & Development
3      33      No      Travel_Frequently      1392      Research & Development
4      27      No      Travel_Rarely      591      Research & Development
...      ...      ...      ...      ...
1465    36      No      Travel_Frequently      884      Research & Development
1466    39      No      Travel_Rarely      613      Research & Development
1467    27      No      Travel_Rarely      155      Research & Development
1468    49      No      Travel_Frequently      1023      Sales
1469    34      No      Travel_Rarely      628      Research & Development

```

```
DistanceFromHome  Education  EducationField  EmployeeCount \
```

0	1	2	Life Sciences	1
1	8	1	Life Sciences	1
2	2	2	Other	1
3	3	4	Life Sciences	1
4	2	1	Medical	1
...	...	...	...	...
1465	23	2	Medical	1
1466	6	1	Medical	1
1467	4	3	Life Sciences	1
1468	2	3	Medical	1
1469	8	3	Medical	1

	EmployeeNumber	...	RelationshipSatisfaction	StandardHours	\
0	1	...	1	80	
1	2	...	4	80	
2	4	...	2	80	
3	5	...	3	80	
4	7	...	4	80	
...	...	...	...	...	
1465	2061	...	3	80	
1466	2062	...	1	80	
1467	2064	...	2	80	
1468	2065	...	4	80	
1469	2068	...	1	80	

	StockOptionLevel	TotalWorkingYears	TrainingTimesLastYear	\
0	0	8	0	
1	1	10	3	
2	0	7	3	
3	0	8	3	
4	1	6	3	
...	...	...	...	
1465	1	17	3	
1466	1	9	5	
1467	1	6	0	
1468	0	17	3	
1469	0	6	3	

	WorkLifeBalance	YearsAtCompany	YearsInCurrentRole	\
0	1	6	4	
1	3	10	7	
2	3	0	0	
3	3	8	7	
4	3	2	2	
...	...	...	...	
1465	3	5	2	
1466	3	7	7	

1467	3	6	2
1468	2	9	6
1469	4	4	3

	YearsSinceLastPromotion	YearsWithCurrManager
0	0	5
1	1	7
2	0	0
3	3	0
4	2	2
...	...	...
1465	0	3
1466	1	7
1467	0	3
1468	0	8
1469	1	2

[1470 rows x 35 columns]>

```
[24]: emp.isnull().sum()
```

```
[24]: Age 0
Attrition 0
BusinessTravel 0
DailyRate 0
Department 0
DistanceFromHome 0
Education 0
EducationField 0
EmployeeCount 0
EmployeeNumber 0
EnvironmentSatisfaction 0
Gender 0
HourlyRate 0
JobInvolvement 0
JobLevel 0
JobRole 0
JobSatisfaction 0
MaritalStatus 0
MonthlyIncome 0
MonthlyRate 0
NumCompaniesWorked 0
Over18 0
OverTime 0
PercentSalaryHike 0
PerformanceRating 0
RelationshipSatisfaction 0
```

```

StandardHours      0
StockOptionLevel    0
TotalWorkingYears   0
TrainingTimesLastYear 0
WorkLifeBalance     0
YearsAtCompany      0
YearsInCurrentRole  0
YearsSinceLastPromotion 0
YearsWithCurrManager 0
dtype: int64

```

### 0.0.3 Step2. [Extract X and y ].

Create X and y columns from the dataframe

```
[25]: X = emp.drop(['Attrition'],axis=1)
      y = emp.Attrition
```

```
[26]: y = y.apply(lambda x:1 if x == 'Yes' else 0)
```

```
[27]: emp.select_dtypes(include=['object']).dtypes
```

```
[27]: Attrition      object
      BusinessTravel  object
      Department     object
      EducationField  object
      Gender          object
      JobRole         object
      MaritalStatus   object
      Over18          object
      OverTime        object
      dtype: object

```

### 0.0.4 Step3. [Feature Engineering]

There are 8 categorical columns (where dtype="object"). Perform one hot encoding and create new columns

```
[28]: emp=pd.
      ↳get_dummies(emp,columns=["BusinessTravel","Department",'EducationField',"Gender","JobRole",
      emp.head()
```

```
[28]:   Age  Attrition  DailyRate  DistanceFromHome  Education  EmployeeCount  \
0   41         Yes       1102                1         2         1
1   49         No        279                8         1         1
2   37         Yes      1373                2         2         1
3   33         No      1392                3         4         1
4   27         No       591                2         1         1

      EmployeeNumber  EnvironmentSatisfaction  HourlyRate  JobInvolvement  ...  \

```

0	1	2	94	3	...
1	2	3	61	2	...
2	4	4	92	2	...
3	5	4	56	3	...
4	7	1	40	3	...

	JobRole_Research Director	JobRole_Research Scientist	\
0	0	0	
1	0	1	
2	0	0	
3	0	1	
4	0	0	

	JobRole_Sales Executive	JobRole_Sales Representative	\
0	1	0	
1	0	0	
2	0	0	
3	0	0	
4	0	0	

	MaritalStatus_Divorced	MaritalStatus_Married	MaritalStatus_Single	\
0	0	0	1	
1	0	1	0	
2	0	0	1	
3	0	1	0	
4	0	1	0	

	Over18_Y	OverTime_No	OverTime_Yes
0	1	0	1
1	1	1	0
2	1	0	1
3	1	0	1
4	1	1	0

[5 rows x 56 columns]

#### 0.0.5 Step4. Now, check shape of X and y.

```
[29]: X = emp.drop(['Attrition'],axis=1)
      X.shape
```

```
[29]: (1470, 55)
```

```
[30]: y.shape
```

```
[30]: (1470,)
```



## Split X and y for training and testing

Create RandomForestClassifier model, fit (no need to scale) and predict

```
[33]: rfc.fit(X_train,y_train)
```

[illegible]

### Print accuracy score between y\_test and y\_pred

```
[35]: 0.8639455782312925
```

Print classification report between y\_test and y\_pred and observe the results

```
[36]: print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.86	0.99	0.92	371
1	0.86	0.17	0.29	70
accuracy			0.86	441
macro avg	0.86	0.58	0.61	441
weighted avg	0.86	0.86	0.82	441

### 0.0.8 Step7. [Feature importance value]

You can look at feature importance values using the property, `rf.feature_importances_`

```
[37]: print(rfc.feature_importances_)
```

```
[6.98321450e-02 3.65227658e-02 2.00159132e-02 5.21457627e-03
0.00000000e+00 1.93136701e-02 3.09477993e-02 2.51635390e-02
1.66715720e-02 5.04884918e-02 1.37024559e-02 1.15212766e-01
1.89715997e-02 1.83688986e-02 1.16183893e-02 6.52783762e-04
9.05995065e-03 0.00000000e+00 2.83057741e-02 6.73115246e-02
5.96985891e-03 1.88628396e-02 5.02462199e-02 1.64563675e-02
8.78154018e-03 4.70126629e-02 3.99898213e-03 1.62801031e-02
2.47672036e-03 5.60666617e-04 3.92298011e-03 4.67197125e-03
2.85590037e-03 1.92092323e-03 3.48874178e-03 2.75324633e-03
3.67258786e-04 4.33772973e-03 1.65455825e-03 1.66961888e-03
3.66392947e-04 7.05642476e-04 4.35437163e-03 3.37746500e-04
1.17707045e-03 6.70465871e-05 4.98737913e-03 6.52221544e-03
1.25986442e-02 2.39045147e-03 3.31770313e-03 2.24531725e-02
0.00000000e+00 9.21168370e-02 9.29418217e-02]
```

Print feature name and its `rf.feature_importances_` values and understand important features

```
[38]: fea_imp = pd.DataFrame(rfc.feature_importances_,index=X_train.
    ↳columns,columns=['Important score']).sort_values(by='Important_
    ↳score',ascending=False)
```

```
[39]: from operator import itemgetter

x = fea_imp.index
y = fea_imp['Important score']

lst = []
print("Feature Name - Feature Importance Score")
print("-----")
```

```

for i in range(55):
    lst.append((x[i], y[i]))
sorted(lst, key=itemgetter(1))
FIS = lst[:]
FIS

```

Feature Name - Feature Importance Score

-----

```

[39]: [('MonthlyIncome', 0.11521276580234263),
      ('OverTime_Yes', 0.09294182174617778),
      ('OverTime_No', 0.09211683703047299),
      ('Age', 0.06983214503413517),
      ('TotalWorkingYears', 0.06731152457797149),
      ('JobLevel', 0.050488491822907974),
      ('YearsAtCompany', 0.05024621988832041),
      ('YearsWithCurrManager', 0.04701266289517826),
      ('DailyRate', 0.036522765839450674),
      ('EnvironmentSatisfaction', 0.030947799266825053),
      ('StockOptionLevel', 0.028305774120963006),
      ('HourlyRate', 0.025163538961806283),
      ('MaritalStatus_Single', 0.022453172508509013),
      ('DistanceFromHome', 0.02001591318605209),
      ('EmployeeNumber', 0.019313670070972264),
      ('MonthlyRate', 0.01897159972501488),
      ('WorkLifeBalance', 0.018862839595341577),
      ('NumCompaniesWorked', 0.018368898619670174),
      ('JobInvolvement', 0.016671571984204086),
      ('YearsInCurrentRole', 0.016456367461815606),
      ('BusinessTravel_Travel_Frequently', 0.016280103058574972),
      ('JobSatisfaction', 0.013702455921311134),
      ('JobRole_Sales Representative', 0.012598644206356679),
      ('PercentSalaryHike', 0.011618389303011392),
      ('RelationshipSatisfaction', 0.009059950651190586),
      ('YearsSinceLastPromotion', 0.008781540180973284),
      ('JobRole_Sales Executive', 0.006522215444916321),
      ('TrainingTimesLastYear', 0.0059698589079211235),
      ('Education', 0.005214576269184216),
      ('JobRole_Research Scientist', 0.004987379130407648),
      ('Department_Sales', 0.004671971253536818),
      ('JobRole_Laboratory Technician', 0.004354371633245582),
      ('EducationField_Technical Degree', 0.00433772972987835),
      ('BusinessTravel_Non-Travel', 0.003998982125104187),
      ('Department_Research & Development', 0.003922980113577673),
      ('EducationField_Marketing', 0.003488741780169301),
      ('MaritalStatus_Married', 0.003317703126868009),
      ('EducationField_Human Resources', 0.002855900366990764),


```

```
( 'EducationField_Medical', 0.0027532463334835272),
( 'BusinessTravel_Travel_Rarely', 0.002476720364451602),
( 'MaritalStatus_Divorced', 0.0023904514748599582),
( 'EducationField_Life Sciences', 0.0019209232322679954),
( 'Gender_Male', 0.0016696188807234103),
( 'Gender_Female', 0.0016545582499910018),
( 'JobRole_Manufacturing Director', 0.0011770704472686203),
( 'JobRole_Human Resources', 0.0007056424757231417),
( 'PerformanceRating', 0.0006527837623514303),
( 'Department_Human Resources', 0.0005606666173878862),
( 'EducationField_Other', 0.00036725878636433484),
( 'JobRole_Healthcare Representative', 0.0003663929467656666),
( 'JobRole_Manager', 0.0003377464999189519),
( 'JobRole_Research Director', 6.704658709318701e-05),
( 'Over18_Y', 0.0),
( 'StandardHours', 0.0),
( 'EmployeeCount', 0.0)]
```

Show a Bar plot between feature column names and feature\_importances\_ score.

```
[40]: pd.Series(rfc.feature_importances_, index=X_train.columns).
      ↪sort_values(ascending=False).plot(kind='bar', figsize=(18,6))
```

[40]: <AxesSubplot:>



## 0.0.9 Step8. [Visualize your RF Decision Tree using graphviz]

<http://www.webgraphviz.com/>.

```
[41]: estim = rfc.estimators_[5]
```

```
[42]: from sklearn import tree
from sklearn.tree import export_graphviz
with open("RFDT.dot", 'w') as f:
    f = tree.export_graphviz(estim, out_file=f, max_depth=4, impurity=False,
        ↳feature_names=X_train.columns.values, class_names=['yes', 'no'],
        ↳rounded=True, proportion=False, precision=2, filled= True)
```

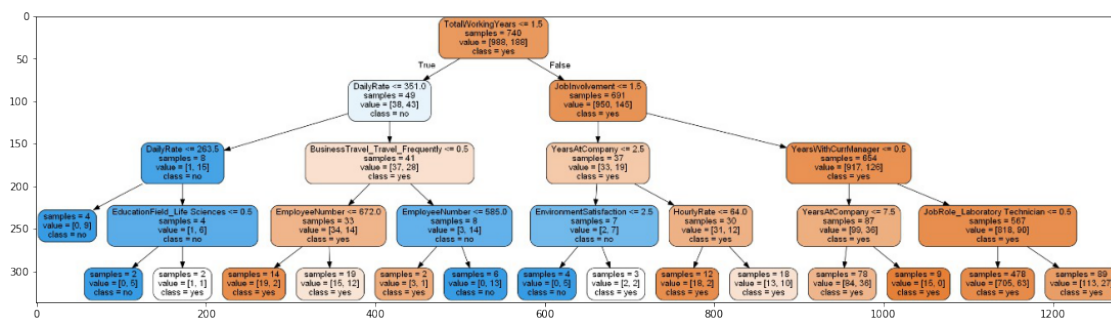
```
[43]: !dot -Tjpg RFDT.dot -o RF.jpg
```

'dot' is not recognized as an internal or external command,  
operable program or batch file.

```
[44]: import matplotlib.pyplot as plt
```

```
image = plt.imread('RF.jpg')
plt.figure(figsize=(19,15))
plt.imshow(image)
```

```
[44]: <matplotlib.image.AxesImage at 0x22fd5fe2f70>
```



### 0.0.10 Step9. [RF with a range of trees]

Fit random forest models with a range of tree numbers [15, 20, 30, 40, 50, 100, 150, 200, 300, 400] and print Out-Of-Bag error for each of these model. Use `model.oob_score__` to get score and subtract this score from 1 to get the oob-error. That is, `oob-error = 1 - model.oob_score__`. Hint: since the only thing changing is the number of trees, the `warm_start` flag can be used so that the model just adds more trees to the existing model each time. Use the `set_params` method to update the number of trees. The following code may help to understand this part.

```
[45]: rf3 =
    ↳RandomForestClassifier(oob_score=True,random_state=42,warm_start=True,n_jobs=-1)
oob_list = list()
# Iterate through all of the possibilities for number of trees
```

```

for n_trees in [15, 20, 30, 40, 50, 100, 150, 200, 300, 400]:
    rf3.set_params(n_estimators=n_trees)
    rf3.fit(X_train, y_train)

    # Get the oob error

    oob_error = 1 - rf3.oob_score_
    oob_list.append(pd.Series({'n_trees': n_trees, 'oob': oob_error}))

rf_oob_df = pd.concat(oob_list, axis=1).T.set_index('n_trees')
rf_oob_df

```

```

[45]:          oob
n_trees
15.0      0.172012
20.0      0.163265
30.0      0.152575
40.0      0.151603
50.0      0.145773
100.0     0.147716
150.0     0.143829
200.0     0.149660
300.0     0.148688
400.0     0.149660

```

### 0.0.11 Step10. [Plot oob-error for each tree]

The following lines will help you

```


[46]: ax = rf_oob_df.plot(legend=False, marker='o', figsize=(10,5))
      ax.set(ylabel='out-of-bag error')

```

```

[46]: [Text(0, 0.5, 'out-of-bag error')]

```



### 0.0.12 Step11. [Compare with DecisionTreeClassifier]

## Create DecisionTreeClassifier, fit and predict on test set

```
[47]: clf2 = DecisionTreeClassifier(criterion='gini',max_depth=4, random_state=42)
```

```
[48]: clf2.fit(X_train,y_train)
```

```
[48]: DecisionTreeClassifier(max_depth=4, random_state=42)
```

```
[49]: y_predict=clf2.predict(X_test)
      y_predict
```

[illegible]

```
0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0,
0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1,
0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0], dtype=int64)
```

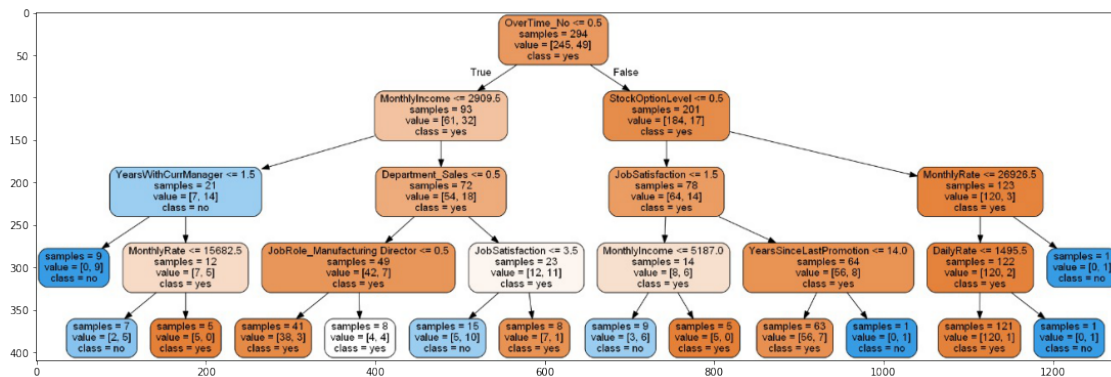
Visualize the tree using graphviz

```
[50]: !dot -Tjpg DTC2.dot -o DT.jpg
```

'dot' is not recognized as an internal or external command,  
operable program or batch file.

```
[51]: image = plt.imread('DT.jpg')
plt.figure(figsize=(19,15))
plt.imshow(image)
```

```
[51]: <matplotlib.image.AxesImage at 0x22fd5d310d0>
```



<http://www.webgraphviz.com/>.

Print accuracy score

```
[52]: accuracy_score(y_test,y_predict)
```

```
[52]: 0.8480725623582767
```

Print classification report

```
[53]: print(classification_report(y_test,y_predict))
```

	precision	recall	f1-score	support
0	0.89	0.94	0.91	371
1	0.53	0.39	0.45	70
accuracy			0.85	441



macro avg	0.71	0.66	0.68	441
weighted avg	0.83	0.85	0.84	441

What is the result of the comparison between RF and DT models? Which gives best accuracy?.

What is your comment on precision, recall, f1 score values?

```
[54]: print("RF model:      ",accuracy_score(y_test,y_pred))
      print("RF Precision: ",precision_score(y_test,y_pred))
      print("RF Recall:    ",recall_score(y_test,y_pred))
      print("RF F1 score:   ",f1_score(y_test,y_pred))
      print("\n")
      print("DT model:      ",accuracy_score(y_test,y_predict))
      print("DT Precision:  ",precision_score(y_test,y_predict))
      print("DT Recall:    ",recall_score(y_test,y_predict))
      print("DT F1 score:   ",f1_score(y_test,y_predict))
```

```
RF model:      0.8639455782312925
RF Precision:  0.8571428571428571
RF Recall:    0.17142857142857143
RF F1 score:   0.2857142857142857
```

```
DT model:      0.8480725623582767
DT Precision:  0.5294117647058824
DT Recall:    0.38571428571428573
DT F1 score:   0.4462809917355372
```